

Elements of Data Science - S20

Midterm Review

This is intended as a guide and is not guaranteed to be inclusive.

Material considered fair for the exam is anything from class and slides.

Note sheet allowed: 1 8x11in sheet, 2-sided

Data Science Tools

- Data Science workflow
- Jupyter+Ipython Notebooks
- conda Virtual Environments
- uses for Git

Python Intro/Review Numpy and Pandas

- Importing modules
- Defining functions
- String Formatting
- What are Exceptions?
- Using assert
- Basic Python data types
- Collections module: Counter, defaultdict
- Python flow control: if: elif: else: , for x in xs:
- Sorting with lambda functions
- List Comprehensions
- Numpy
 - arrays
 - indexing/slicing
 - Boolean masks and bitwise operations
- Pandas
 - Series
 - DataFrames
 - indexing/slicing
 - .describe

Visualization and Data Exploration

- Matplotlib
 - plotting using matplotlib
 - plt vs ax level plotting
 - using plt.subplots()
- Variable Types
- Central tendencies
 - mean
 - median
- Spread
 - variance

- std deviation
- skew
- IQR
- Correlation
 - Pearson Correlation Coefficient
- Univariate Plotting
 - hist
 - seaborn.distplot()
 - boxplots
- Bivariate Plotting
 - scatter
 - jointplot
 - pairplot
- Categorical Plotting
 - bar
 - catplot

Hypothesis Testing

- Random Sampling vs Population Distribution
- Sample Statistic
- Confidence Intervals
- Normal (Gaussian) Distribution
 - Standard Normal Distribution
 - Z-Score
- Central Limit Theorem
- Bootstrap Sampling
- A/B Test
- Hypothesis Testing
 - Type I and II error
 - Significance and Power
 - Permutation Tests
 - One-tailed vs Two-tailed
 - p-values
- Calculating “How many observations?”
 - what 4 values are related?
- Multi-Armed Bandit
 - benefits of using
 - greedy
 - epsilon-greedy

Modeling, Prediction Model Evaluation and Selection

- Dimensions of ML
 - Interpretation vs Prediction
 - Learning Paradigms (SL,UL,etc.)
 - Regression vs Classification
 - Binary, Multiclass, Multilabel
- sklearn common functions
 - .fit
 - .predict
 - .predict_proba

- .score
- Generalization
 - Train/Test split
 - stratification
- Overfitting/Underfitting
 - Bias/Variance Tradeoff
- Baseline Models
- Tuning Hyperparameters and Model Selection
 - k-Fold Cross Validation
 - Grid Search
- Plotting Model Fit
 - Validation Curve
 - Learning Curve
- Metrics: Classification
 - Confusion Matrix
 - Accuracy/Error
 - Precision
 - Recall
 - Precision-Recall Curve
 - F1 Score
 - ROC Curve (FPR vs TPR)
 - ROC AUC
- Metrics: Regression
 - R^2
 - Adjusted R^2
 - Mean Squared Error
 - RMSE

Machine Learning Models

- Concept of Gradient Descent
- k-NearestNeighbor
- Naive Bayes (generally)
- Simple Linear Regression
 - Residuals in linear models
 - Interpreting Coefficients of OLS
 - Colinearity
- Multiple Linear Regression
- Logistic Regression
 - Interpreting Coefficients of LogReg
- Decision Trees
- Ensembles
 - Random Forest
 - Gradient Boost
 - Stacking